



## Configuração do Hive 3.0 e Cloudera HUE em um ecossistema Hadoop!

Criado em Agosto/2018

Autor: Vinícius José Leodido Correia

linkedin: <https://www.linkedin.com/in/vinileodido/>

Primeiramente iremos instalar o Hive 3.0 e em seguida o HUE, um Self Service BI mantido pela Cloudera sobre licença Apache com uma UX bastante significativa.

Porém, antes de instalarmos o Hive alguns requisitos são necessários.

- Postgresql

Iremos instalar os clients do Postgresql e o Postgresql Server.

```
[hduser@hdpmaster ~]$ sudo yum list postgresql postgresql-server
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
* base: mirror.globo.com
* epel: mirror.globo.com
* extras: mirror.globo.com
* updates: mirror.globo.com
Available Packages
postgresql.i686                               9.2.23-3.el7_4                                base
postgresql.x86_64                             9.2.23-3.el7_4                                base
postgresql-server.x86_64                     9.2.23-3.el7_4                                base
[hduser@hdpmaster ~]$
```

Instalando o Postgres



```
[hduser@hdpmaster ~]$ sudo yum install postgresql postgresql-server -y
```

```
Loaded plugins: fastestmirror
```

```
Loading mirror speeds from cached hostfile
```

```
* base: mirror.globo.com
```

```
* epel: mirror.globo.com
```

```
* extras: mirror.globo.com
```

```
* updates: mirror.globo.com
```

```
Resolving Dependencies
```

```
--> Running transaction check
```

```
---> Package postgresql.x86_64 0:9.2.23-3.el7_4 will be installed
```

```
--> Processing Dependency: postgresql-libs(x86-64) = 9.2.23-3.el7_4 for package: postgresql-9.2.23-3.el7_4.x86_64
```

```
--> Processing Dependency: libpq.so.5()(64bit) for package: postgresql-9.2.23-3.el7_4.x86_64
```

```
---> Package postgresql-server.x86_64 0:9.2.23-3.el7_4 will be installed
```

```
--> Running transaction check
```

```
---> Package postgresql-libs.x86_64 0:9.2.23-3.el7_4 will be installed
```

```
--> Finished Dependency Resolution
```

```
Dependencies Resolved
```

Package	Arch	Version	Repository	Size
Installing:				
postgresql	x86_64	9.2.23-3.el7_4	base	3.0 M
postgresql-server	x86_64	9.2.23-3.el7_4	base	3.8 M
Installing for dependencies:				
postgresql-libs	x86_64	9.2.23-3.el7_4	base	234 k

```
Transaction Summary
```

```
Install 2 Packages (+1 Dependent package)
```

```
Total download size: 7.0 M
```

```
Installed size: 33 M
```

```
Downloading packages:
```

(1/3): postgresql-libs-9.2.23-3.el7_4.x86_64.rpm	234 kB	00:00:00
(2/3): postgresql-9.2.23-3.el7_4.x86_64.rpm	3.0 MB	00:00:00
(3/3): postgresql-server-9.2.23-3.el7_4.x86_64.rpm	3.8 MB	00:00:00

```
Total
```

```
6.0 MB/s | 7.0 MB 00:00:01
```

```
Running transaction check
```

```
Running transaction test
```

```
Transaction test succeeded
```

```
Running transaction
```

Installing : postgresql-libs-9.2.23-3.el7_4.x86_64	1/3
Installing : postgresql-9.2.23-3.el7_4.x86_64	2/3
Installing : postgresql-server-9.2.23-3.el7_4.x86_64	3/3
Verifying : postgresql-9.2.23-3.el7_4.x86_64	1/3
Verifying : postgresql-server-9.2.23-3.el7_4.x86_64	2/3
Verifying : postgresql-libs-9.2.23-3.el7_4.x86_64	3/3



```
Installed:
  postgresql.x86_64 0:9.2.23-3.el7_4                postgresql-server.x86_64 0:9.2.23-3.el7_4

Dependency Installed:
  postgresql-libs.x86_64 0:9.2.23-3.el7_4

Complete!
[hduser@hdpmaster ~]$
```

Inicialmente, está desabilitado no sistema:

```
[hduser@hdpmaster ~]$ sudo systemctl status postgresql
• postgresql.service - PostgreSQL database server
  Loaded: loaded (/usr/lib/systemd/system/postgresql.service; disabled; vendor preset: disabled)
  Active: inactive (dead)
```

Iniciando o postgres:

```
hduser@hdpmaster:~
[hduser@hdpmaster ~]$ sudo systemctl enable postgresql
Created symlink from /etc/systemd/system/multi-user.target.wants/postgresql.service to /usr/lib/systemd/system/postgresql.service.
[hduser@hdpmaster ~]$ sudo postgresql-setup initdb
Initializing database ... /var/lib/pgsql/data is not writeable by postgres
[hduser@hdpmaster ~]$ sudo chown -R postgres:postgres /var/lib/pgsql/data
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$ sudo postgresql-setup initdb
Initializing database ... OK

[hduser@hdpmaster ~]$ sudo systemctl start postgresql
[hduser@hdpmaster ~]$ sudo systemctl status postgresql
• postgresql.service - PostgreSQL database server
  Loaded: loaded (/usr/lib/systemd/system/postgresql.service; enabled; vendor preset: disabled)
  Active: active (running) since Mon 2018-08-06 09:02:00 -03; 13s ago
  Process: 1746 ExecStart=/usr/bin/pg_ctl start -D ${PGDATA} -s -o -p ${PGPORT} -w -t 300 (code=exited, status=0/SUCCESS)
  Process: 1741 ExecStartPre=/usr/bin/postgresql-check-db-dir ${PGDATA} (code=exited, status=0/SUCCESS)
  Main PID: 1749 (postgres)
  CGroup: /system.slice/postgresql.service
          └─1749 /usr/bin/postgres -D /var/lib/pgsql/data -p 5432
             └─1750 postgres: logger process
                └─1752 postgres: checkpoint process
                   └─1753 postgres: writer process
                      └─1754 postgres: wal writer process
                         └─1755 postgres: autovacuum launcher process
                            └─1756 postgres: stats collector process

Aug 06 09:01:59 hdpmaster systemd[1]: Starting PostgreSQL database server...
Aug 06 09:02:00 hdpmaster systemd[1]: Started PostgreSQL database server.
[hduser@hdpmaster ~]$
```

Efetue o download do Hive (Version 3.0.0) através de um dos mirrors do endereço <https://hive.apache.org/downloads.html> e descompacte na pasta /usr/local; em seguida iremos criar o symlink para o diretório /opt/hive:



```
[hduser@hdpmaster ~]$ wget http://ftp.unicamp.br/pub/apache/hive/hive-3.0.0/apache-hive-3.0.0-bin.tar.gz
--2018-08-06 09:13:47-- http://ftp.unicamp.br/pub/apache/hive/hive-3.0.0/apache-hive-3.0.0-bin.tar.gz
Resolving ftp.unicamp.br (ftp.unicamp.br)... 143.106.10.149
Connecting to ftp.unicamp.br (ftp.unicamp.br)|143.106.10.149|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 307672459 (293M) [application/x-gzip]
Saving to: 'apache-hive-3.0.0-bin.tar.gz'

100%[=====>] 307,672,459 5.86MB/s in 58s

2018-08-06 09:14:45 (5.06 MB/s) - 'apache-hive-3.0.0-bin.tar.gz' saved [307672459/307672459]

[hduser@hdpmaster ~]$ sudo tar xzf apache-hive-3.0.0-bin.tar.gz -C /usr/local
[hduser@hdpmaster ~]$ cd /usr/local/
[hduser@hdpmaster local]$ ls -l
total 0
drwxr-xr-x. 10 root root 184 Aug 6 09:15 apache-hive-3.0.0-bin
drwxr-xr-x. 2 root root 6 Apr 11 01:59 bin
drwxr-xr-x. 2 root root 6 Apr 11 01:59 etc
drwxr-xr-x. 2 root root 6 Apr 11 01:59 games
drwxr-xr-x. 11 hduser hadoop 172 Aug 5 17:40 hadoop-3.1.0
drwxr-xr-x. 2 root root 6 Apr 11 01:59 include
drwxr-xr-x. 7 10 143 245 Jul 7 05:09 jdk1.8.0_181
drwxr-xr-x. 2 root root 6 Apr 11 01:59 lib
drwxr-xr-x. 2 root root 6 Apr 11 01:59 lib64
drwxr-xr-x. 2 root root 6 Apr 11 01:59 libexec
drwxr-xr-x. 2 root root 6 Apr 11 01:59 sbin
drwxr-xr-x. 5 root root 49 Aug 5 07:11 share
drwxr-xr-x. 2 root root 6 Apr 11 01:59 src
[hduser@hdpmaster local]$ sudo mv apache-hive-3.0.0-bin hive-3.0.0
[hduser@hdpmaster local]$ sudo ln -s /usr/local/hive-3.0.0/ /opt/hive
[hduser@hdpmaster local]$
```

Edite as variáveis de ambiente para o Hive:

```
$ vi ~/.bash_profile

#Adicione as seguintes variaveis abaixo
#HIVE_HOME
export HIVE_HOME=/opt/hive
export PATH=$PATH:$HIVE_HOME/bin
export CLASSPATH=$CLASSPATH:$HADOOP_HOME/lib/*:.
export CLASSPATH=$HIVE_HOME/lib/*:.
```



```
# .bash_profile

# Get the aliases and functions
if [ -f ~/.bashrc ]; then
    . ~/.bashrc
fi

# User specific environment and startup programs

PATH=$PATH:$HOME/.local/bin:$HOME/bin

export PATH

#JAVA_HOME
export JAVA_HOME=/opt/jdk
export PATH=$PATH:$JAVA_HOME/bin

#HADOOP_HOME
export HADOOP_HOME=/opt/hadoop
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin

#HIVE_HOME
export HIVE_HOME=/opt/hive
export PATH=$PATH:$HIVE_HOME/bin
export CLASSPATH=$CLASSPATH:$HADOOP_HOME/lib/*:.
export CLASSPATH=$HIVE_HOME/lib/*:.
```

Para verificar se as configurações foram provisionadas corretamente no sistema verifique com o seguinte comando:

```
$ hive --version
```

```
[hduser@hdpmaster ~]$ source ~/.bash_profile
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$ hive --version
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive-3.0.0/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-3.1.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive 3.0.0
Git git://vgargwork.local/Users/vgarg/repos/hive.apache.master.latest -r ce61711a5fa54ab34fc74d86d521ecaeea6b072a
Compiled by vgarg on Fri May 18 11:38:33 PDT 2018
From source with checksum 81fcb93b608965ed7ac968bae1187fab
[hduser@hdpmaster ~]$
```

Antes de iniciarmos os serviços do Hive, precisaremos utilizar o script de criação do Metastore (Banco de dados que armazenará os metadados do Hive) no Postgresql;

Por que estamos fazendo este procedimento do Metastore para o Postgresql? O Metastore vem preparado para utilização através do Derby, um pequeno sistema de banco de dados local que é utilizado pelo Hive, porém, o mesmo só permitirá que uma sessão ativa seja instanciada para uso do serviço. Quando efetuamos a migração do Derby para um RDBMS como Oracle, MySQL e Postgres, conseguimos utilizar mais de uma sessão simultaneamente.

Faça a instalação do driver jdbc Postgresql para a pasta de bibliotecas do Hive:



```
[hduser@hdpmaster ~]$ wget https://jdbc.postgresql.org/download/postgresql-42.2.4.jar
--2018-08-06 10:12:46-- https://jdbc.postgresql.org/download/postgresql-42.2.4.jar
Resolving jdbc.postgresql.org (jdbc.postgresql.org)... 174.143.35.228, 2001:4800:1501:1::228
Connecting to jdbc.postgresql.org (jdbc.postgresql.org)|174.143.35.228|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 814992 (796K) [application/java-archive]
Saving to: 'postgresql-42.2.4.jar'

100%[=====>] 814,992      738KB/s   in 1.1s

2018-08-06 10:12:48 (738 KB/s) - 'postgresql-42.2.4.jar' saved [814992/814992]

[hduser@hdpmaster ~]$ sudo cp postgresql-42.2.4.jar /opt/hive/lib/
[hduser@hdpmaster ~]$
```

Localize o diretório de scripts Sql (Schema versão 3.0.0) do Metastore para o Postgres no diretório de instalação do Hive:

```
[hduser@hdpmaster ~]$ ls -l /opt/hive/scripts/metastore/upgrade/postgres/hive*schema*3.0.0*
-rw-r--r--. 1 root root 52500 May 15 18:42 /opt/hive/scripts/metastore/upgrade/postgres/hive-schema-3.0.0.postgres.sql
[hduser@hdpmaster ~]$
```

Crie o Metastore Database e a conta de usuario no Postgres:

```
$ sudo -u postgres psql

postgres=# CREATE USER hiveuser WITH PASSWORD 'mypassword';
postgres=# CREATE DATABASE metastore;
postgres=# GRANT ALL PRIVILEGES ON DATABASE metastore to hiveuser;
postgres=# \q
```

```
[hduser@hdpmaster ~]$ sudo -u postgres psql
could not change directory to "/home/hduser"
psql (9.2.23)
Type "help" for help.

postgres=# CREATE USER hiveuser WITH PASSWORD 'h1v3u$3r';
CREATE ROLE
postgres=# CREATE DATABASE metastore;
CREATE DATABASE
postgres=# \du

              List of roles
Role name | Attributes | Member of
-----+-----+-----
hiveuser |           | {}
postgres | Superuser, Create role, Create DB, Replication | {}

postgres=# GRANT ALL PRIVILEGES ON DATABASE metastore to hiveuser;
GRANT
postgres=# \q
[hduser@hdpmaster ~]$
```

Vamos editar o arquivo pg\_hba.conf para que possamos acessar o host com os usuários que criaremos daqui em diante:



```
[hduser@hdpmaster ~]$ sudo ls -l /var/lib/pgsql/data/pg_hba.conf
[sudo] password for hduser:
-rw-----. 1 postgres postgres 4232 Aug  6 09:01 /var/lib/pgsql/data/pg_hba.conf
[hduser@hdpmaster ~]$ sudo vi /var/lib/pgsql/data/pg_hba.conf
```

```
#
# This file is read on server startup and when the postmaster receives
# a SIGHUP signal.  If you edit the file on a running system, you have
# to SIGHUP the postmaster for the changes to take effect.  You can
# use "pg_ctl reload" to do that.
#
# Put your actual configuration here
# -----
#
# If you want to allow non-local connections, you need to add more
# "host" records.  In that case you will also need to make PostgreSQL
# listen on a non-local interface via the listen_addresses
# configuration parameter, or via the -i or -h command line switches.


# TYPE  DATABASE        USER            ADDRESS                 METHOD

# "local" is for Unix domain socket connections only
local   all         all                                peer
# IPv4 local connections:
host    all         all             127.0.0.1/32            ident
host    all         all             0.0.0.0/0               md5
# IPv6 local connections:
host    all         all             ::1/128                 ident
# Allow replication connections from localhost, by a user with the
# replication privilege.
#local   replication  postgres        peer
#host    replication  postgres        127.0.0.1/32            ident
#host    replication  postgres        ::1/128                 ident
~
~
~
~
```

Em seguida iremos editar o arquivo postgresql.conf :

```
#-----
# CONNECTIONS AND AUTHENTICATION
#-----

# - Connection Settings -

listen_addresses = '*'
#listen_addresses = 'localhost'      # what IP address(es) to listen on;
                                     # comma-separated list of addresses;
                                     # defaults to 'localhost'; use '*' for all
                                     # (change requires restart)
#port = 5432                          # (change requires restart)
```

Reinicie o postgresql:

```
$ sudo systemctl restart postgresql
```

Entre novamente no postgresql com o usuario postgres e altere para que o database a ser utilizado seja o metastore; Em seguida, execute o script de criação do hive-schema-3.0.0.postgres.sql conforme comando abaixo:



```
$ sudo -u postgres psql

postgres=# \c metastore
metastore=# \i /usr/local/hive-3.0.0/scripts/metastore/upgrade/postgres/hive-schema-3.0.0.postgres.sql
```

```
[hduser@hdpmaster ~]$ sudo -u postgres psql
could not change directory to "/home/hduser"
psql (9.2.23)
Type "help" for help.

postgres=# \c metastore
You are now connected to database "metastore" as user "postgres".
metastore=# \i /usr/local/hive-3.0.0/scripts/metastore/upgrade/postgres/hive-schema-3.0.0.postgres.sql █
```

Agora precisamos conceder permissão para todas os objetos do metastore para o usuário que utilizaremos "hiveuser":

```
$ sudo -u postgres psql

postgres=# \c metastore
metastore=# \set tuples_only on
metastore=# \o /tmp/grant-privs
metastore=# SELECT 'GRANT SELECT, INSERT, UPDATE, DELETE ON "' || schemaname || '".' || tablename ||'" TO hiveuser ;'
metastore=# FROM pg_tables
metastore=# WHERE tableowner = CURRENT_USER and schemaname = 'public';
metastore=# \o
metastore=# \set tuples_only off
metastore=# \i /tmp/grant-privs
```

```
[root@hdpmaster ~]# sudo -u postgres psql
could not change directory to "/root"
psql (9.2.23)
Type "help" for help.

postgres=# \c metastore
You are now connected to database "metastore" as user "postgres".
metastore=# \set tuples_only on
Showing only tuples.
metastore=# \o /tmp/grant-privs
metastore=# SELECT 'GRANT SELECT, INSERT, UPDATE, DELETE ON "' || schemaname || '".' || tablename ||'" TO hiveuser ;'
metastore=# FROM pg_tables
metastore=# WHERE tableowner = CURRENT_USER and schemaname = 'public';
metastore=# \o
metastore=# \set tuples_only off
Tuples only is off.
metastore=# \i /tmp/grant-privs █
```

Desconecte do postgresql:





```
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
GRANT
metastore=# \q
[root@hdpmaster ~]#
```

Vamos editar o arquivo hive-site.xml :

```
[hduser@hdpmaster ~]$ cd /opt/hive/conf/
[hduser@hdpmaster conf]$ ls -l
total 328
-rw-r--r--. 1 root root 1596 May 15 18:42 beeline-log4j2.properties.template
-rw-r--r--. 1 root root 298702 May 18 15:38 hive-default.xml.template
-rw-r--r--. 1 root root 2365 May 15 18:42 hive-env.sh.template
-rw-r--r--. 1 root root 2274 May 15 18:42 hive-exec-log4j2.properties.template
-rw-r--r--. 1 root root 3086 May 15 18:42 hive-log4j2.properties.template
-rw-r--r--. 1 root root 2060 May 15 18:42 ivysettings.xml
-rw-r--r--. 1 root root 3558 May 15 18:42 llap-cli-log4j2.properties.template
-rw-r--r--. 1 root root 7163 May 15 18:42 llap-daemon-log4j2.properties.template
-rw-r--r--. 1 root root 2662 May 15 18:42 parquet-logging.properties
[hduser@hdpmaster conf]$ vi hive-site.xml
```

```
<?xml version="1.0" encoding="UTF-8" standalone="no"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>!--
Licensed to the Apache Software Foundation (ASF) under one or more
contributor license agreements. See the NOTICE file distributed with
this work for additional information regarding copyright ownership.
The ASF licenses this file to You under the Apache License, Version 2.0
(the "License"); you may not use this file except in compliance with
the License. You may obtain a copy of the License at
```

<http://www.apache.org/licenses/LICENSE-2.0>

```
Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.
-->
```

```
<configuration>
  <property>
    <name>javax.jdo.option.ConnectionURL</name>
    <value>jdbc:postgresql://192.168.1.200:5432/metastore</value>
  </property>

  <property>
    <name>javax.jdo.option.ConnectionDriverName</name>
    <value>org.postgresql.Driver</value>
  </property>

  <property>
```



```
<name>javax.jdo.option.ConnectionUserName</name>
<value>hiveuser</value>
</property>

<property>
<name>javax.jdo.option.ConnectionPassword</name>
<value>h1v3m3t4$t0r3</value>
</property>

<property>
<name>hive.server2.thrift.port</name>
<value>10000</value>
<description>HiveServer2 thrift interface - TCP port number to listen on, default 10000</description>
</property>

<property>
<name>hive.server2.thrift.bind.host</name>
<value>localhost</value>
<description>HiveServer2 bind host</description>
</property>

<property>
<name>datanucleus.autoCreateSchema</name>
<value>>false</value>
</property>

<property>
<name>hive.metastore.schema.validation</name>
<value>true</value>
</property>
</configuration>
```

Verificando se o Hive está funcionando:

```
[hduser@hdpmaster ~]$ hive --version
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive-3.0.0/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-3.1.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive 3.0.0
Git git://vgargwork.local/Users/vgarg/repos/hive.apache.master.latest -r ce61711a5fa54ab34fc74d86d521ecaeea6b072a
Compiled by vgarg on Fri May 18 11:38:33 PDT 2018
From source with checksum 81fcb93b608965ed7ac968bae1187fab
[hduser@hdpmaster ~]$
```

Iniciando o schematool do Hive (Atenção! Como executamos o sql de criação do Hive manualmente, alguns Warnings poderão ser exibidos em tela!):



```
[hduser@hdpmaster ~]$ ls -l $HIVE_HOME/bin
total 44
-rwxr-xr-x. 1 root root 881 May 15 18:42 beeline
drwxr-xr-x. 3 root root 4096 Aug 6 09:15 ext
-rwxr-xr-x. 1 root root 10026 May 15 18:42 hive
-rwxr-xr-x. 1 root root 1900 May 15 18:42 hive-config.sh
-rwxr-xr-x. 1 root root 885 May 15 18:42 hiveserver2
-rwxr-xr-x. 1 root root 880 May 15 18:42 hplsql
-rwxr-xr-x. 1 root root 3064 May 15 18:42 init-hive-dfs.sh
-rwxr-xr-x. 1 root root 832 May 15 18:42 metatool
-rwxr-xr-x. 1 root root 884 May 15 18:42 schematool
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$ schematool -dbType postgres -initSchema
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive-3.0.0/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-3.1.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Metastore connection URL: jdbc:postgresql://192.168.1.200:5432/metastore
Metastore Connection Driver : org.postgresql.Driver
Metastore connection User: hiveuser
Starting metastore schema initialization to 3.0.0
Initialization script hive-schema-3.0.0.postgres.sql

Initialization script completed
schemaTool completed
[hduser@hdpmaster ~]$
```

Iremos instanciar o serviço do HiveServer2:

Importante averiguar nos arquivos de configuração do Hive, se as propriedades do Thrift foram adicionadas ao arquivo hive-site.xml:

```
<property>
  <name>hive.server2.thrift.port</name>
  <value>10000</value>
  <description>TCP port number to listen on, default 10000</description>
</property>
```



```
[hduser@hdpmaster ~]$ ls -l $HIVE_HOME/bin
total 44
-rwxr-xr-x. 1 root root 881 May 15 18:42 beeline
drwxr-xr-x. 3 root root 4096 Aug 6 09:15 ext
-rwxr-xr-x. 1 root root 10026 May 15 18:42 hive
-rwxr-xr-x. 1 root root 1900 May 15 18:42 hive-config.sh
-rwxr-xr-x. 1 root root 885 May 15 18:42 hiveserver2
-rwxr-xr-x. 1 root root 880 May 15 18:42 hplsql
-rwxr-xr-x. 1 root root 3064 May 15 18:42 init-hive-dfs.sh
-rwxr-xr-x. 1 root root 832 May 15 18:42 metatool
-rwxr-xr-x. 1 root root 884 May 15 18:42 schematool
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$ nohup hive --service hiveserver2 &
[1] 4092
[hduser@hdpmaster ~]$ nohup: ignoring input and appending output to 'nohup.out'

[hduser@hdpmaster ~]$
```

Vamos verificar se o hiveserver2 está em execução no background do sistema:

```
[hduser@hdpmaster ~]$ jps
4247 Jps
1928 SecondaryNameNode
1708 NameNode
4092 RunJar
2142 ResourceManager
[hduser@hdpmaster ~]$ ps -aef | grep 4092
hduser 4092 3665 28 16:10 pts/0 00:00:12 /opt/jdk/bin/java -Dproc_jar -Dproc_hiveserver2 -Dlog4j.config
line.UnsupportedTerminal -Dyarn.log.dir=/opt/hadoop/logs -Dyarn.log.file=hadoop.log -Dyarn.home.dir=/opt/hadoop
p.log.file=hadoop.log -Dhadoop.home.dir=/opt/hadoop -Dhadoop.id.str=hduser -Dhadoop.root.logger=INFO,console -i
ve-service-3.0.0.jar org.apache.hive.service.server.HiveServer2
hduser 4259 3665 0 16:11 pts/0 00:00:00 grep --color=auto 4092
[hduser@hdpmaster ~]$
```

Antes de testarmos o beeline, será necessário a alteração de algumas propriedades no arquivo core-site.xml nos arquivos de configuração do hadoop:

```
$ vi $HADOOP_CONF_DIR/core-site.xml
```

```
<property>
  <name>hadoop.proxyuser.<user_connected_to_hdfs>.hosts</name>
  <value>*</value>
</property>
<property>
  <name>hadoop.proxyuser.<user_connected_to_hdfs>.groups</name>
  <value>*</value>
</property>

<property>
  <name>hadoop.proxyuser.hduser.hosts</name>
  <value>*</value>
</property>
<property>
  <name>hadoop.proxyuser.hduser.groups</name>
  <value>*</value>
</property>
```

```
#Interrompa os serviços do Hadoop
$ stop-dfs.sh; stop-yarn.sh
```



```
#Reinicie novamente os serviços  
$ start-dfs.sh; start-yarn.sh
```

Vamos testar o beeline:

```
$ beeline  
  
beeline> !connect jdbc:hive2://hdpmaster:10000 scott tiger org.apache.hive.jdbc.HiveDriver
```

```
[hduser@hdpmaster ~]$ beeline  
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/usr/local/hive-3.0.0/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/usr/local/hadoop-3.1.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]  
Beeline version 3.0.0 by Apache Hive  
beeline> !connect jdbc:hive2://hdpmaster:10000 scott tiger org.apache.hive.jdbc.HiveDriver  
Connecting to jdbc:hive2://hdpmaster:10000  
Connected to: Apache Hive (version 3.0.0)  
Driver: Hive JDBC (version 3.0.0)  
Transaction isolation: TRANSACTION_REPEATABLE_READ  
0: jdbc:hive2://hdpmaster:10000> show tables;  
INFO : Compiling command(queryId=hduser_20180806162937_975f84c3-50e4-4052-8a09-a74b70fdccec): show tables  
INFO : Concurrency mode is disabled, not creating a lock manager  
INFO : Semantic Analysis Completed  
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:tab_name, type:string, comment:from deserializer)], properties:null)  
INFO : Completed compiling command(queryId=hduser_20180806162937_975f84c3-50e4-4052-8a09-a74b70fdccec); Time taken: 0.8 seconds  
INFO : Concurrency mode is disabled, not creating a lock manager  
INFO : Executing command(queryId=hduser_20180806162937_975f84c3-50e4-4052-8a09-a74b70fdccec): show tables  
INFO : Starting task [Stage-0:DDL] in serial mode  
INFO : Completed executing command(queryId=hduser_20180806162937_975f84c3-50e4-4052-8a09-a74b70fdccec); Time taken: 0.038 seconds  
INFO : OK  
INFO : Concurrency mode is disabled, not creating a lock manager  
+-----+  
| tab_name |  
+-----+  
+-----+  
No rows selected (1.154 seconds)  
0: jdbc:hive2://hdpmaster:10000> !quit  
Closing: 0: jdbc:hive2://hdpmaster:10000  
[hduser@hdpmaster ~]$
```

Instalando as dependencias do HUE

como usuário root, instale os seguintes pacotes

```
$ yum install ant  
$ yum install gcc g++  
$ yum install make  
$ yum install python-devel.x86_64  
$ yum groupinstall "Development Tools"  
$ yum install krb5-devel  
$ yum install libxslt-devel libxml2-devel  
$ yum install mysql-devel.x86_64  
  
$ yum install ncurses-devel zlib-devel texinfo gtk+-devel gtk2-devel  
$ yum install qt-devel tcl-devel tk-devel kernel-headers kernel-devel
```



```
$ yum install yum install gmp-devel.x86_64
$ yum install sqlite-devel.x86_64
$ yum install cyrus-sasl.x86_64

$ yum install postfix system-switch-mail cyrus-imapd cyrus-plain cyrus-md5 cyrus-utilsum install postfix system-switch-mail cyrus-imapd cyrus-plain cyrus-md5 cyrus-utils

$ yum install memcached-devel.x86_64
$ yum install libevent libevent-devel
$ yum install postfix
$ yum install cyrus-sasl
$ yum install cyrus-imapd
$ yum install openldap-devel
```

#### Atenção:

Como o pacote ant instala o openJdk e realiza um "bypass" na instalação que você efetuou no sistema e aprovisionou nas variáveis de ambiente, liste as variáveis simbólicas que estão no diretório /etc/alternatives/jav\* em um arquivo de bkp e em seguida, as remova do diretório;

```
[hduser@hdpmaster ~]$ sudo ls -l /etc/alternatives/jav* >> alternatives_java.txt
[hduser@hdpmaster ~]$ cat alternatives_java.txt
lrwxrwxrwx. 1 root root 73 Aug 6 17:13 /etc/alternatives/java -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64/jre/bin/java
lrwxrwxrwx. 1 root root 77 Aug 6 17:13 /etc/alternatives/java.1.gz -> /usr/share/man/man1/java-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 70 Aug 6 17:13 /etc/alternatives/javac -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64/bin/javac
lrwxrwxrwx. 1 root root 78 Aug 6 17:13 /etc/alternatives/javac.1.gz -> /usr/share/man/man1/javac-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 72 Aug 6 17:13 /etc/alternatives/javadoc -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64/bin/javadoc
lrwxrwxrwx. 1 root root 80 Aug 6 17:13 /etc/alternatives/javadoc.1.gz -> /usr/share/man/man1/javadoc-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 70 Aug 6 17:13 /etc/alternatives/javah -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64/bin/javah
lrwxrwxrwx. 1 root root 78 Aug 6 17:13 /etc/alternatives/javah.1.gz -> /usr/share/man/man1/javah-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 70 Aug 6 17:13 /etc/alternatives/javap -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64/bin/javap
lrwxrwxrwx. 1 root root 78 Aug 6 17:13 /etc/alternatives/javap.1.gz -> /usr/share/man/man1/javap-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 60 Aug 6 17:13 /etc/alternatives/java_sdk -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64
lrwxrwxrwx. 1 root root 60 Aug 6 17:13 /etc/alternatives/java_sdk.1.gz -> /usr/share/man/man1/java_sdk-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 68 Aug 6 17:13 /etc/alternatives/java_sdk.1.gz -> /usr/share/man/man1/java_sdk-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 60 Aug 6 17:13 /etc/alternatives/java_sdk_openjdk -> /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64
lrwxrwxrwx. 1 root root 68 Aug 6 17:13 /etc/alternatives/java_sdk_openjdk.1.gz -> /usr/share/man/man1/java_sdk_openjdk-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 68 Aug 6 17:13 /etc/alternatives/java_sdk_openjdk.1.gz -> /usr/share/man/man1/java_sdk_openjdk-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
lrwxrwxrwx. 1 root root 68 Aug 6 17:13 /etc/alternatives/java_sdk_openjdk.1.gz -> /usr/share/man/man1/java_sdk_openjdk-java-1.8.0-openjdk-1.8.0.181-3.b13.el7_5.x86_64.1.gz
[hduser@hdpmaster ~]$ sudo rm /etc/alternatives/jav*
[hduser@hdpmaster ~]$ sudo ls -l /etc/alternatives/jav*
ls: cannot access /etc/alternatives/jav*: No such file or directory
[hduser@hdpmaster ~]$
```

Execute novamente um source em suas variáveis de ambiente (Para que o sistema confirma novamente qual o Jdk está aprovisionado):

```
$ source ~/.bash_profile
```

Por fim, execute o comando java -version a fim de verificar qual o Jdk realmente seu sistema estará utilizando:

```
$ java -version
```

```
[hduser@hdpmaster ~]$ java -version
java version "1.8.0_181"
Java(TM) SE Runtime Environment (build 1.8.0_181-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.181-b13, mixed mode)
[hduser@hdpmaster ~]$
```

Em seguida, realize a instalação do Maven em seu sistema:

<https://maven.apache.org/download.cgi>

```
$ wget http://mirror.nbtelcom.com.br/apache/maven/maven-3/3.5.4/binaries/apache-maven-3.5.4-bin.tar.gz
```



```
$ sudo tar xzf apache-maven-3.5.4-bin.tar.gz -C /usr/local

$ sudo mv apache-maven-3.5.4-bin maven-3.5.4
$ sudo ln -s /usr/local/maven-3.5.4/opt/maven

$ vi ~/.bash_profile

#MAVEN_HOME
export MAVEN_HOME=/opt/maven
export PATH=$PATH:$MAVEN_HOME/bin
```

Para averiguar se a instalação ocorreu com sucesso:

```
$ mvn -version
```

```
[hduser@hdpmaster ~]$ mvn -version
Apache Maven 3.5.4 (1edded0938998edf8bf061f1ceb3cfdeccf443fe; 2018-06-17T15:33:14-03:00)
Maven home: /opt/maven
Java version: 1.8.0_181, vendor: Oracle Corporation, runtime: /usr/local/jdk1.8.0_181/jre
Default locale: en_US, platform encoding: UTF-8
OS name: "linux", version: "3.10.0-862.9.1.el7.x86_64", arch: "amd64", family: "unix"
[hduser@hdpmaster ~]$
```

Continuaremos instalando mais alguns pacotes necessários:

```
$ sudo yum install asciidoc cyrus-sasl-devel cyrus-sasl-gssapi cyrus-sasl-plain
$ sudo yum install krb5-devel libffi-devel libxml2-devel libxslt-devel openldap-devel gmp-devel sqlite-devel
$ sudo yum install python-devel
$ sudo yum install mysql mysql-devel

$ sudo yum install libtidy
$ sudo yum install openssl-devel
```

Por fim, iremos efetuar a instalação do HUE:

Acesse a página da distribuição em <http://gethue.com>

```
$ wget https://www.dropbox.com/s/0rhlrljmyw6bnfc/hue-4.2.0.tgz

$ sudo tar xzf hue-4.2.0.tgz -C /usr/local

$ cd /usr/local
$ pushd hue-4.2.0/
$ sudo make install

# Após a instalação
$ popd
```

Iremos alterar as permissões de acesso para a pasta que foi criada:



```
[hduser@hdpmaster local]$ ls -lah
total 0
drwxr-xr-x. 18 root   root   236 Aug  6 18:05 .
drwxr-xr-x. 13 root   root   155 Aug  5 07:11 ..
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 bin
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 etc
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 games
drwxr-xr-x. 11 hduser  hadoop 172 Aug  5 17:40 hadoop-3.1.0
drwxr-xr-x. 10 root   root   184 Aug  6 09:15 hive-3.0.0
drwxr-xr-x.  7 root   root   242 Aug  6 18:08 hue
drwxrwxr-x.  9 hduser  hadoop 244 Aug  6 18:05 hue-4.2.0
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 include
drwxr-xr-x.  7      10    143 245 Jul  7 05:09 jdk1.8.0_181
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 lib
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 lib64
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 libexec
drwxr-xr-x.  6 root   root    99 Aug  6 17:28 maven-3.5.4
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 sbin
drwxr-xr-x.  5 root   root    49 Aug  5 07:11 share
drwxr-xr-x.  2 root   root    6 Apr 11 01:59 src
[hduser@hdpmaster local]$
```

```
$ sudo useradd hue
$ sudo chown -R hue:hue /usr/local/hue
```

Remova a pasta que utilizamos para realizar o build do HUE:

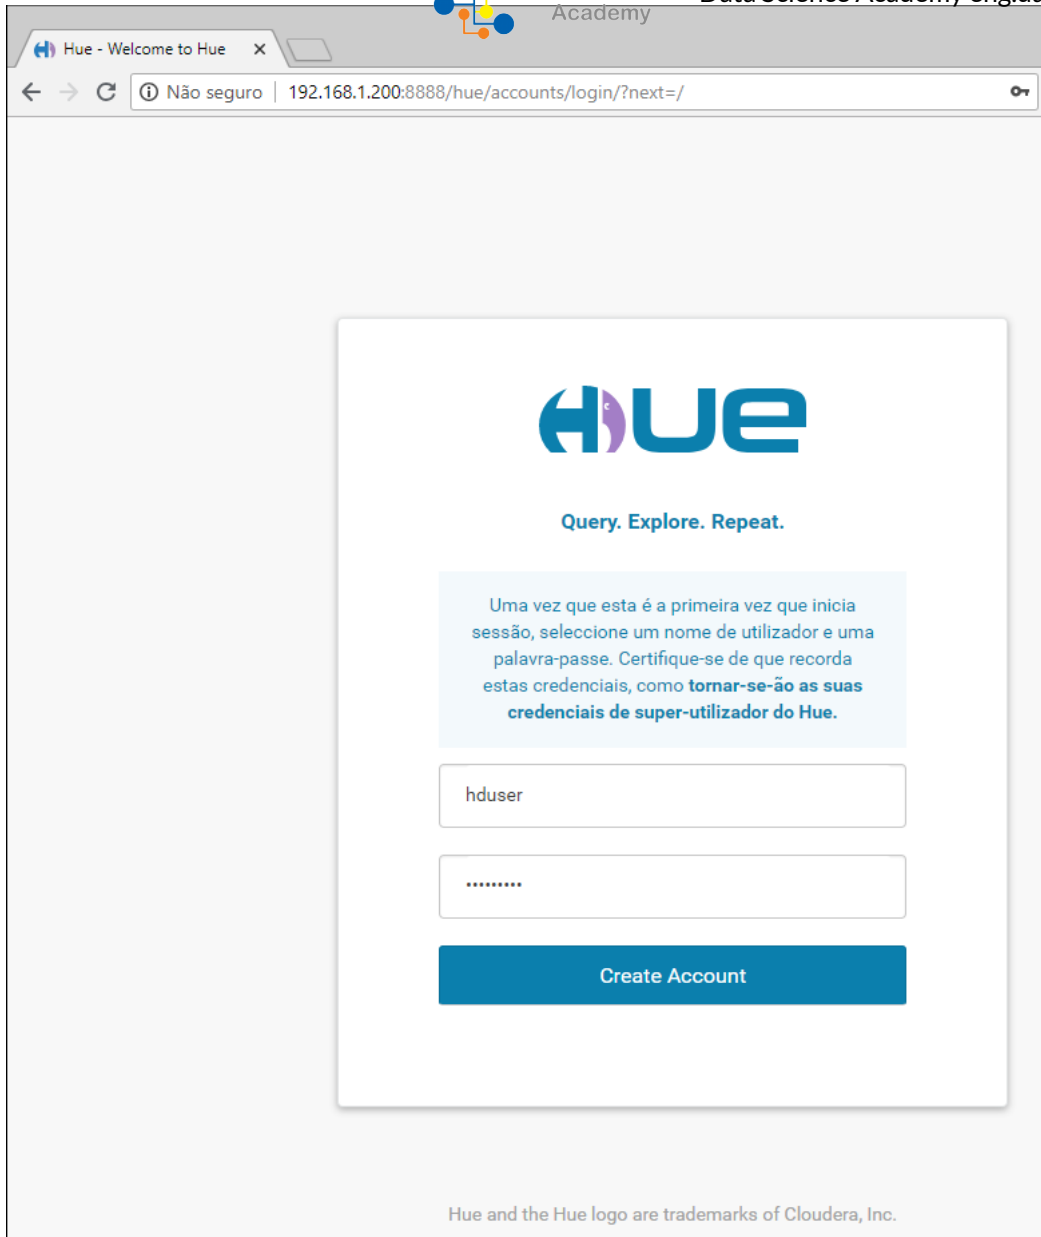
```
$ sudo rm -r /usr/local/hue-4.2.0
$ sudo ln -s /usr/local/hue /opt/hue
$ vi ~/.bash_profile
```

Iremos inicializar o HUE:

```
$ sudo $HUE_HOME/build/env/bin/supervisor
```

Acesse o browser de sua preferência, com o endereço do Host onde está instalado o HUE utilizando a porta :8888





Efetue Logout do HUE, e vamos para o próximo passo onde iremos migrar do Sqlite para um RDBMS.

Interrompa o serviço do HUE, com o seguinte comando:

```
$ ps -f -u hue
```



```
#Localize o PID do processo runcherrypserver e realize um kill no usuário root  
#Repita por pelo menos 4 vezes.
```

#### Importante!

O Hue é inicializado através do comando supervisor [sudo \$HUE\_HOME/build/env/bin/supervisor], o qual desempenha a função de monitorar se o processo runcherrypserver está em execução e caso o mesmo não esteja localizado na pilha de processos do sistema, o Supervisor instancia um novo processo do runcherrypserver; Porém, se o processo runcherrypserver for identificado "abort" ou "kill" pelo supervisor em um curto intervalo de tempo, por sucedidas vezes (x4) o mesmo irá interromper as tentativas de inicialização do processo runcherrypserver e o HUE terá seus serviços paralisados, até que seja executado novamente o processo do Supervisor ou sejam corrigidos os Logs de Erros (caso seja o cenário).

```
[hduser@hdpmaster ~]$ ps -f -u hue  
UID      PID  PPID  C  STIME TTY          TIME CMD  
hue      31811 31539 12 19:32 pts/0    00:00:01 /usr/local/hue/build/env/bin/python2.7 /usr/local/hue/build/env/bin/hue runcherrypserver  
[hduser@hdpmaster ~]$ sudo kill 31811  
[hduser@hdpmaster ~]$ ps -f -u hue  
UID      PID  PPID  C  STIME TTY          TIME CMD  
[1]+  Exit 1      sudo $HUE_HOME/build/env/bin/supervisor  
[hduser@hdpmaster ~]$
```

#### Configurações do HUE:

```
$HUE_HOME/desktop/conf
```

#### Criando o acesso a um banco para os metadados do HUE:

Assim como vimos anteriormente (Hive utilizando o Derby como BD local para seus metadados), o Hue armazena seus metadados também em um repositório local com o Sqlite; Para finalidades de testes, seu uso é recomendado, porém, para uma abordagem em que serão utilizadas diversas conexões simultâneas e com diversos usuários acessando o HUE, efetuar uma migração para um RDBMS como Oracle, MySQL ou PostgreSQL, se faz necessário, conforme recomendação da Cloudera.

O passo a passo poderá ser encontrado nos seguintes endereços (mas deixei as etapas que realizei descritas em sequência):

<http://cloudera.github.io/hue/latest/admin-manual/manual.html#configuring-to-access-another-database>

<https://community.hortonworks.com/questions/399/how-to-move-hue-database-from-default-sqlite-datab.html>

```
$ sudo -u postgres psql
```

```
postgres=# CREATE USER hueuser WITH password 'hu3u$3r';  
postgres=# CREATE DATABASE hue;  
postgres=# GRANT ALL PRIVILEGES ON DATABASE hue TO hueuser;
```

Crie um diretório de backup; Efetue o dump dos dados do HUE para este diretório através do seguinte comando:

```
$ sudo $HUE_HOME/build/env/bin/hue dumpdata > ~/bkp_conf/hue_db_dump.json
```

Efetue também, um backup do arquivo de configurações do HUE antes de efetuarmos modificações no mesmo:

```
$ sudo cp $HUE_HOME/desktop/conf/hue.ini ~/bkp_conf/
```

Abra o arquivo hue.ini no diretório de configurações do HUE [\$HUE\_HOME/desktop/conf/hue.ini] e altere as seguintes propriedades:

```
$ sudo vi $HUE_HOME/desktop/conf/hue.ini
```

```
#Localize os parametros de configurações [[database]] e altere informando as respectivas credenciais de seu usuário
```

```
engine=postgresql_psycopg2  
host=192.168.1.200 #Host do PostgreSQL  
port=5432  
user=hueuser  
password=hu3u$3r  
name=hue
```



```
schema=public
```

Por fim, as alterações estarão conforme abaixo

```
[[database]]
# Database engine is typically one of:
# postgresql_psycopg2, mysql, sqlite3 or oracle.
#
# Note that for sqlite3, 'name', below is a path to the filename. For other backends, it is the database name
# Note for Oracle, options={"threaded":true} must be set in order to avoid crashes.
# Note for Oracle, you can use the Oracle Service Name by setting "host=" and "port=" and then "name=<host>:<port>/<service_name>".
# Note for MariaDB use the 'mysql' engine.
## engine=sqlite3
engine=postgresql_psycopg2
host=192.168.1.200
port=5432
user=hueuser
password=hu3u$3r
# conn_max_age option to make database connection persistent value in seconds
# https://docs.djangoproject.com/en/1.9/ref/databases/#persistent-connections
## conn_max_age=0
# Execute this script to produce the database password. This will be used when 'password' is not set.
## password_script=/path/script
## name=desktop/desktop.db
name=hue
## options={}
# Database schema, to be used only when public schema is revoked in postgres
## schema=public

# Configuration options for specifying the Desktop session.
# For more info, see https://docs.djangoproject.com/en/1.4/topics/http/sessions/
```

Importante! Se realmente for utilizar o Postgresql como repositório para os metadados do HUE, a instalação da biblioteca psycopg2 do Python se faz necessário pois a base do HUE é um servidor web em Django!

```
$ sudo pip install psycopg2
$ sudo pip install psycopg2-binary
```

```
[hduser@hdpmaster ~]$ sudo pip install psycopg2
Collecting psycopg2
  Downloading https://files.pythonhosted.org/packages/7c/e6/d5161798a5e8900f24216cb730f2c2be5e4758a80d35c8588306831c0c99/psycopg2-2.7.5-cp27-cp27mu-manylinux1\_x86\_64.whl (2.7MB)
    100% |#####| 2.7MB 1.7MB/s
Installing collected packages: psycopg2
Successfully installed psycopg2-2.7.5
[hduser@hdpmaster ~]$
```

Execute os seguintes comandos abaixo (Importante que o HUE esteja com os serviços interrompidos!)

```
$ sudo $HUE_HOME/build/env/bin/hue syncdb --noinput
$ sudo $HUE_HOME/build/env/bin/hue migrate
$ sudo $HUE_HOME/build/env/bin/hue loaddata ~/bkp_conf/hue_db_dump.json

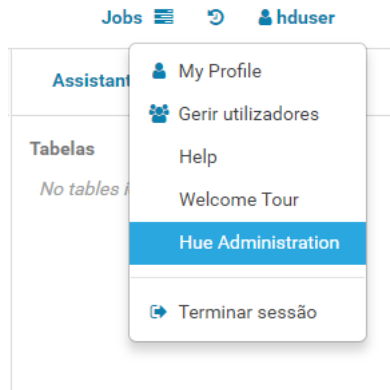
#Inicie novamente o Hue pelo supervisor
$ sudo nohup $HUE_HOME/build/env/bin/supervisor &
```

Pronto, efetuamos a migração do sqlite para o Postgresql!

Ao efetuar acesso novamente no HUE com o seu usuário cadastrado, exibirão diversas mensagens de erros; Isto ocorre pois o ambiente não está devidamente mapeado de acordo com as configurações que vieram em default no arquivo \$HUE\_HOME/desktop/conf/hue.ini



Para conferir as mensagens de Log, acesso a opção "HUE Administration" bem próximo do canto superior direito, onde estão as informações do usuário que está logado no momento.



Conforme a tela abaixo, vemos quais respectivas variáveis necessitam de alteração / atualização no arquivo de configuração:



## Assistente de início rápido - Hue 4.2.0 - Query. Explore. Repeat.

Passo 1: Verificar  
configuração

Passo 2: Exemplos

Passo 3: Utilizadores

Passo 4: Ir!

### A verificar configuração actual

Ficheiros de configuração localizados em `/usr/local/hue/desktop/conf`

Possíveis configurações incorrectas detectadas. Repare e reinicie o Hue.

<code>hadoop.hdfs_clusters.default.webhdfs_url</code>	Valor actual: <code>http://localhost:50070/webhdfs/v1</code> Não foi possível aceder à raiz do sistema de ficheiros
<code>Resource Manager</code>	Failed to contact an active Resource Manager: YARN RM returned a failed response: HTTPConnectionPool(host='localhost', port=8088): Max retries exceeded with url: /ws/v1/cluster/apps?user=hduser&user.name=hue&doAs=hduser (Caused by NewConnectionError(: Failed to establish a new connection: [Errno 111] Connection refused'))
<code>Hive</code>	Failed to access Hive warehouse: /user/hive/warehouse
<code>HBase Browser</code>	The application won't work without a running HBase Thrift Server v1.
<code>Impala</code>	No available Impalad to send queries to.
<code>Oozie Editor/Dashboard</code>	A aplicação não funciona sem um servidor Oozie em execução
<code>Pig Editor</code>	A aplicação não funciona sem um servidor Oozie em execução
<code>Hadoop Security: Sentry Service</code>	Failed to connect to Sentry API (version 1).
<code>Hadoop Security: Sentry Service</code>	Failed to connect to Sentry API (version 2).
<code>Spark</code>	The app won't work without a running Livy Spark Server

## Throubleshooting nas configurações do HUE

### 1. Habilitando a restapi webhdfs

Para corrigirmos o primeiro Warning, necessitaremos atuar diretamente nos arquivos de configuração do Hadoop em seu ecossistema.

Localize o arquivo `hdfs-site.xml` em seu diretório de configuração do Hadoop e acrescente a seguinte propriedade:

```
<property>
  <name>dfs.webhdfs.enabled</name>
  <value>true</value>
</property>
```

Reinicie os serviços do seu HDFS;

Acesse ao arquivo de configuração do HUE `[$HUE_HOME/desktop/conf/hue.ini]` e localize a seguinte tag `[[hdfs_clusters]]`

Altere para a porta correta do seu NameNode (Estou utilizando a versão 3.1.0, portanto a porta está situada na :9870 )

```
[[hdfs_clusters]]
# HA support by using HttpFs

[[[default]]]
# Enter the filesystem uri
fs_defaults=hdfs://localhost:8020
```



```
# NameNode logical name.
## logical_name=

# Use WebHdfs/HttpFs as the communication mechanism.
# Domain should be the NameNode or HttpFs host.
# Default port is 14000 for HttpFs.
## webhdfs_url=http://localhost:50070/webhdfs/v1
webhdfs_url=http://localhost:9870/webhdfs/v1
```

## 2. Alterando para a porta correta do Journal do Yarn [Resource Manager]

Acesse ao arquivo de configuração do HUE [\$HUE\_HOME/desktop/conf/hue.ini] e localize a seguinte tag [[yarn\_clusters]]  
Altere para a porta correta do seu NameNode (Estou utilizando a versão 3.1.0, portanto a porta está situada na :8088 )

```
[[yarn_clusters]]

[[[default]]]
# Enter the host on which you are running the ResourceManager
## resourcemanager_host=localhost

# The port where the ResourceManager IPC listens on
## resourcemanager_port=8032
resourcemanager_port=8088
```

## 3. Alterando o superuser do HDFS nas configurações do HUE

Acesse ao arquivo de configuração do HUE [\$HUE\_HOME/desktop/conf/hue.ini] e localize a seguinte tag [desktop]  
Localize a seguinte informação "default\_hdfs\_superuser=hdfs" e altere para o usuário que está utilizando como superuser HDFS

```
# This should be the hadoop cluster admin
## default_hdfs_superuser=hdfs
default_hdfs_superuser=hduser
```

## 4. Alterando as Configurações do Hive

Acesse ao arquivo de configuração do HUE [\$HUE\_HOME/desktop/conf/hue.ini] e localize a seguinte tag [beeswax]  
Informe qual a localização do arquivo hive-site.xml em seu sistema

```
[beeswax]

# Host where HiveServer2 is running.
# If Kerberos security is enabled, use fully-qualified domain name (FQDN).
## hive_server_host=localhost

# Port where HiveServer2 Thrift server runs on.
## hive_server_port=10000

# Hive configuration directory, where hive-site.xml is located
## hive_conf_dir=/etc/hive/conf
hive_conf_dir=/opt/hive/conf
```

Certifique-se de que os serviços do hiveserver2 estão em execução; Caso contrário, basta instanciar através do comando:

```
$ nohup hive --service hiveserver2 &
```

Interrompa os serviços do HUE e os reinicie novamente;

Efetue acesso com suas credenciais de acesso, e pronto a interface do HUE com Integração ao Hive e HDFS estão funcionais e operando em conjunto com as demais ferramentas de seu ecossistema.



The screenshot shows the Hue web interface for Hive. The browser address bar displays the URL: 192.168.1.200:8888/hue/editor?editor=16&type=hive. The interface includes a top navigation bar with the Hue logo, a search bar for saved documents, and a user profile for 'hduser'. The left sidebar shows the 'Hive' section with a list of databases: 'default' and 'my\_first\_db'. The main editor area shows a Hive SQL query being executed: 

```
1 create database my_first_db;
2
3 use my_first_db;
4
5 create table my_first_hive_table (column_a string, integer_first int);
6
```

 Below the editor, the 'Query History' section displays a list of executed queries with their status and timestamps. The queries are: 

- um minuto atrás: create table my\_first\_hive\_table (column\_a string, integer\_first int)
- um minuto atrás: use my\_first\_db
- 2 minutos atrás: create database my\_first\_db
- 2 minutos atrás: create database my\_first\_db; use my\_first\_db; create table my\_first\_hive\_table (column\_a string, integer\_first int);

Vamos consultar se os diretórios com os respectivos nomes das tabelas foram criados no HDFS:

```
[hduser@hdpmaster ~]$ hdfs dfs -ls /user/hive/warehouse
Found 1 items
drwxr-xr-x - hduser supergroup          0 2018-08-06 20:54 /user/hive/warehouse/my_first_db.db
[hduser@hdpmaster ~]$
[hduser@hdpmaster ~]$ hdfs dfs -ls /user/hive/warehouse/my_first_db.db
Found 1 items
drwxr-xr-x - hduser supergroup          0 2018-08-06 20:54 /user/hive/warehouse/my_first_db.db/my_first_hive_table
[hduser@hdpmaster ~]$
```



Data Science  
Academy

Data Science Academy [eng.davidborges@gmail.com](mailto:eng.davidborges@gmail.com) 59532d8f5e4cdead748b456a

Agora fica ao seu encargo explorar mais da interface e configurações do HUE, assim como habilitar / desabilitar os serviços que não serão utilizados no momento e realizar troubleshooting quando necessário!